

CS 231A Section: Computer Vision Libraries Overview



Amir Sadeghian

Feb 2018

Overview

- **Opencv**
- Deep Learning Frameworks
 - **Caffe**
 - Torch
 - Tensorflow

Other CV libraries

- **Vlfeat:** An Open source library with popular computer vision algorithms specializing in image understanding and local features extraction and matching.
- **scikit-learn:** An open source Python library that implements a range of machine learning, preprocessing, cross-validation and visualization algorithms.
- **PCL:** A standalone, large scale, open project for 2D/3D image and point cloud processing.
- **SLAM frameworks (bundler, visualsfm, meshlab):** Applications for 3D reconstruction using structure from motion (SFM).
- **Libraries for specific tasks:** e.g. tracking libraries, Detection libraries ...

OpenCV

Introduction to OpenCV

- Open source computer vision and machine learning library
- Contains implementations of a large number of vision algorithms
- Written natively in C++, also has C, Python, Java, and MATLAB interfaces
- Supports Windows, Linux, Mac OS X, Android, and iOS

Installation

- Download from <http://opencv.org> and compile from source
- Windows: Run executable downloaded from OpenCV website
- Mac OS X: Install through MacPorts, `easy_install`, ...
- Linux: Install through the package manager (e.g. `yum`, `apt`) but make sure the version is sufficiently up-to-date for your needs

Basic OpenCV Structures

- **Point, Point2f** - 2D Point
- **Size** - 2D size structure
- **Rect** - 2D rectangle object
- **RotatedRect** - Rect object with angle
- **Mat** - image object
- ...

Point

- 2D Point Object

- int x, y;

- Sample Functions

- Point.dot(<Point>) - computes dot product

- Point.inside(<Rect>) - returns true if point is inside

```
10 int main(int argc, char* argv[]){
11     Point a(1,1);
12     Point b(2,2);
13     Point c = a+b;
14
15     cout << c.x << ", " << c.y << endl;
16     c = a*2;
17     cout << c.x << ", " << c.y << endl;
18     cout << norm(b) << endl;
19
20     if(a==b)
21         cout << "A == B" << endl;
22     else
23         cout << "A != B" << endl;
24     if(b != c)
25         cout << "B != C" << endl;
26     else
27         cout << "B == C" << endl;
28
29     return 0;
30 }
```

Math operators,
you may use

- Point operator +

- Point operator +=

- Point operator -

- Point operator -=

- Point operator *

- Point operator *=

- bool operator ==

- bool operator !=

- double norm

Size

- 2D Size Structure
 - int width, height;
- Functions
 - Size.area() - returns (width * height)

Rect

- 2D Rectangle Structure
 - int x, y, width, height;
- Functions
 - Rect.tl() - **return top left point**
 - Rect.br() - **return bottom right point**

CV::Mat

- The primary data structure in OpenCV is the Mat object. It stores images and their components.

- Main items

- rows, cols - length and width(int)
- channels - 1: grayscale, 3: BGR
- depth: CV_<bit_depth>C<#chan>

```
int main(int argc, char* argv[]){
    Mat image = imread(argv[1]);
    cout << "Columns = " << image.cols << endl;
    cout << "Rows   = " << image.rows << endl;
    cout << "Type   = ";
    if(image.type() == CV_8UC1) cout << "CV_8UC1" << endl;
    else if(image.type() == CV_8UC3) cout << "CV_8UC3" << endl;
    else if(image.type() == CV_32FC1) cout << "CV_32FC1" << endl;
    else if(image.type() == CV_32FC3) cout << "CV_32FC3" << endl;
    else cout << "Unknown" << endl;
    return 0;
}
```

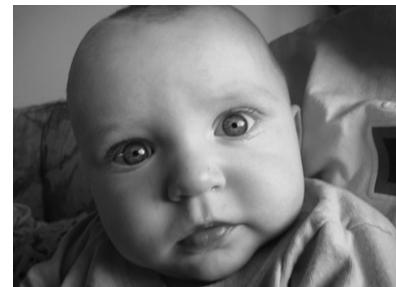
- See the manuals for more information

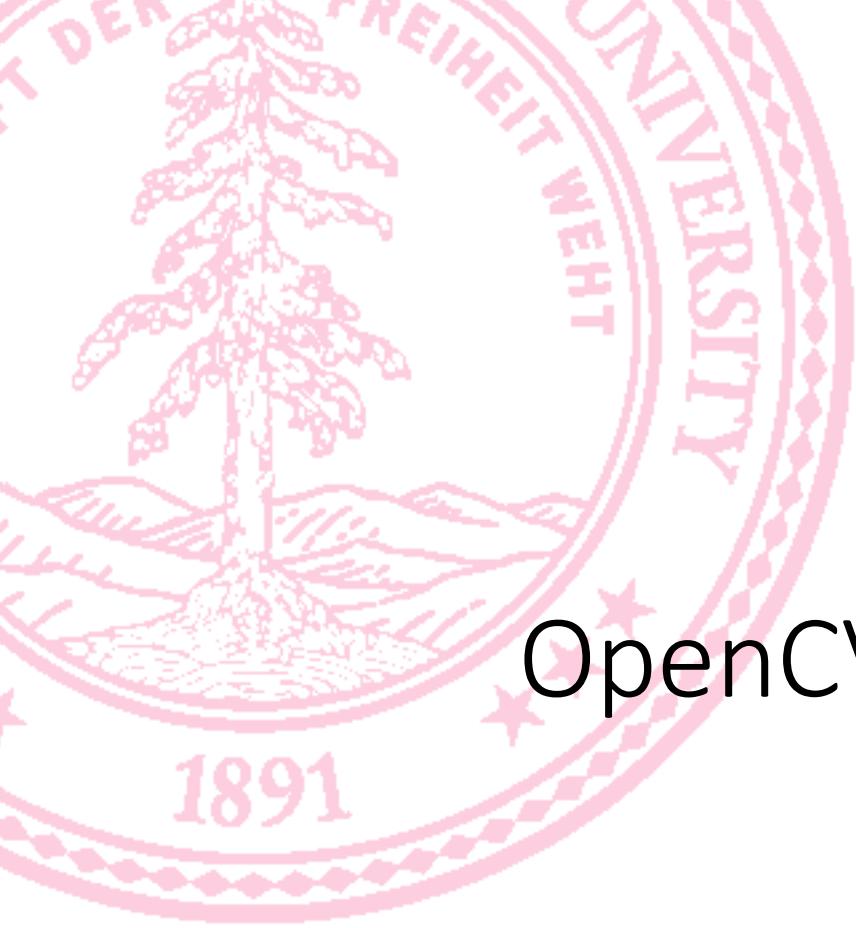
cv::Mat- Functions

- Mat.**at<datatype>**(row, col) - returns pointer to a location in the image
- Mat.**channels()** - returns the number of channels
- Mat.**clone()** - returns a deep copy of the image
- Mat.**create(rows, cols, TYPE)** - re-allocates new memory to matrix
- Mat.**cross(<Mat>)** - computes cross product of two matrices
- Mat.**depth()** - returns bit-depth of a matrix
- Mat.**dot(<Mat>)** - computes the dot product of two matrices

Pixeltypes

- PixelTypes shows how the image is represented in data
 - BGR - The default color of imread(). Normal 3 channel color
 - HSV - Hue is color, Saturation is amount, Value is lightness. 3 channels
 - GRayscale - Gray values, Single channel
- OpenCV requires that images be in BGR or Grayscale in order to be shown or saved. Otherwise, undesirable effects may appear.





OpenCV Functions



Image Normalization and Thresholding

- Normalization remaps a range of pixel values to another range of pixel values
 - `void normalize(InputArray src,OutputArray dst,...)`
- OpenCV provides a general purpose method for thresholding an image
 - `double threshold(InputArray src,OutputArray dst,double thresh,double maxval,int type)`
 - Specify thresholding scheme specified by the type variable

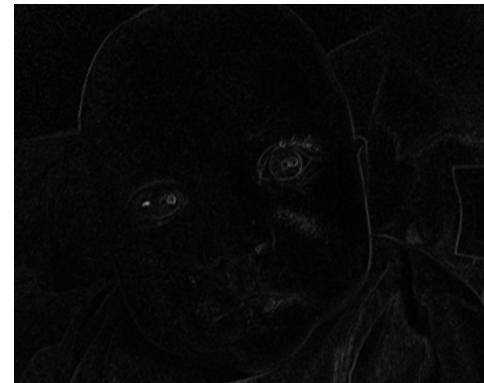


Image Smoothing

- Reduces the sharpness of edges and smooths out details in an image
- OpenCV implements several of the most commonly used methods
 - void GaussianBlur(InputArray src,OutputArray dst,...)
 - void medianBlur(InputArray src,OutputArray dst,...)
- Other functions include generic convolution, separable convolution, dilate, and erode.

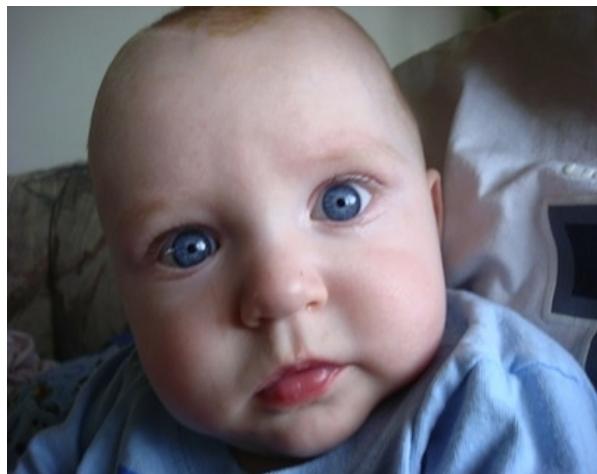
Image Smoothing: Code

```
#include <cv.h>
#include <cvaux.h>
#include <highgui.h>

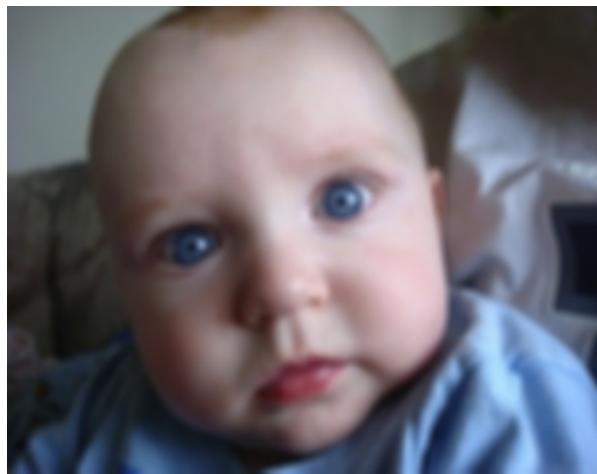
int main(int argc, char** argv) {
    //Read in colored image
    cv::Mat image=cv::imread(argv[1]);
    cv::imwrite("photo.jpg",image);
    //Apply Gaussian blur
    cv::Mat image_gaussian_blur;
    image.convertTo(image_gaussian_blur,CV_8UC3);
    cv::GaussianBlur(image_gaussian_blur,image_gaussian_blur,cv::Size(0,0),9);
    cv::imwrite("photo_gaussian_blur.jpg",image_gaussian_blur);

    //Apply median blur
    cv::Mat image_median_blur;
    image.convertTo(image_median_blur,CV_8UC3);
    cv::medianBlur(image_median_blur,image_median_blur,17);
    cv::imwrite("photo_median_blur.jpg",image_median_blur);
}
```

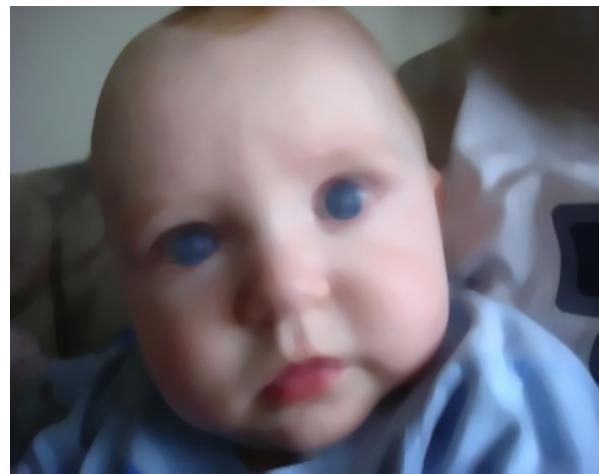
Image Smoothing: Sample Image



Original



Gaussian Blur



Median Blur

Edge Detection

- OpenCV implements a number of operators to help detect edges in an image
 - Sobel Operator
 - `void cv::Sobel(image in, image out, CV_DEPTH, dx, dy);`
 - Scharr Operator
 - `void cv::Scharr(image in, image out, CV_DEPTH, dx, dy);`
 - Laplacian Operator
 - `void cv::Laplacian(image in, image out, CV_DEPTH);`
- OpenCV also implements multi-stage edge detection algorithms such as Canny edge detection
- Tip: If your image is noisy, then edge detection will often exaggerate the noise
- Sometimes smoothing the image before running edge detection gives better results

Edge Detection: Code

```
#include <cv.h>
#include <cvaux.h>
#include <highgui.h>

int main(int argc, char** argv){
    //Read image as grayscale, delete zero to read in color
    cv::Mat image=cv::imread(argv[1],0);
    cv::imwrite("photo_gray.jpg",image);

    //Calculate x-gradient using Sobel operator
    cv::Mat image_gradient_x;
    image.convertTo(image_gradient_x,CV_32FC1);
    cv::Sobel(image_gradient_x,image_gradient_x,CV_32FC1,0,1);
    //Absolute value and normalize
    cv::convertScaleAbs(image_gradient_x,image_gradient_x);

    //Calculate y-gradient using Sobel operator
    cv::Mat image_gradient_y;
    image.convertTo(image_gradient_y,CV_32FC1);
    cv::Sobel(image_gradient_y,image_gradient_y,CV_32FC1,1,0);
    //Absolute value and normalize
    cv::convertScaleAbs(image_gradient_y,image_gradient_y);

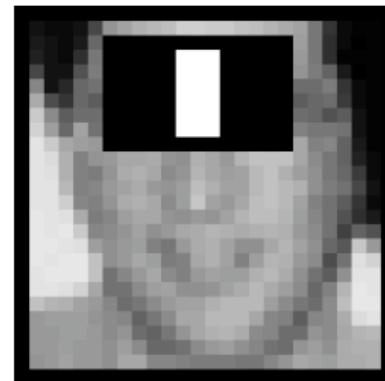
    //Average the x and y gradients into one image
    cv::Mat image_gradient;
    cv::addWeighted(image_gradient_x,0.5,image_gradient_y,0.5,0,image_gradient);
    cv::imwrite("photo_gradient.jpg",image_gradient);
}
```

Edge Detection: Sample Results



Face Detection: Viola-Jones

- ▶ Robust and fast
- ▶ Introduced by Paul Viola and Michael Jones
 - ▶ http://research.microsoft.com/~viola/Pubs/Detect/violaJones_CVPR2001.pdf
- ▶ Haar-like Features



And Many More ...

- [Object Tracking using OpenCV](#)
- [Handwritten Digits Classification : An OpenCV \(C++ / Python \) Tutorial](#)
- [Eye Detector using OpenCV](#)
- [Image Recognition and Object Detection](#)
- [Head Pose Estimation using OpenCV](#)
- [Configuring Qt for OpenCV](#)
- ...



Deep Learning Frameworks



Deep Learning Frameworks

- Caffe
- Torch/PyTorch
 - NYU
 - scientific computing framework in Lua
 - supported by Facebook
- TensorFlow
 - Google
 - Python
- Theano/Pylearn2
 - U. Montreal
 - Python
 - symbolic computation and automatic differentiation
- MatConvNet
 - Oxford U.
 - Deep Learning in MATLAB

Framework Comparison

- More alike than different
 - All express deep models
 - All are open-source (contributions differ)
 - Most include scripting for hacking and prototyping
- No strict winners, experiment and choose the framework that best fits your work

Caffe: Overview

- What is Caffe?
- Training/Finetuning a simple model
- Deep dive into Caffe!

What is Caffe?

- A deep learning framework
- Open framework, models, and worked examples for deep learning
- 4000+ citations, 250+ contributors, 11,000+ forks
- Focus has been vision, but branching out: sequences, reinforcement learning, speech + text

Caffe

- Pure C++ / CUDA architecture for deep learning
 - command line, Python, MATLAB interfaces
- Fast, well-tested code
- Tools, reference models, demos, and recipes
- Switch between CPU and GPU
 - `Caffe::set_mode(Caffe::GPU);`

Installation

- <http://caffe.berkeleyvision.org/installation.html> •
- CUDA, OPENCV
- BLAS (Basic Linear Algebra Subprograms): operations like matrix multiplication, matrix addition, both implementation for CPU(cBLAS) and GPU(cuBLAS). provided by MKL(INTEL), ATLAS, openBLAS, etc.
- Boost: a c++ library. > Use some of its math functions and shared_pointer.
- glog,gflags provide logging & command line utilities. > Essential for debugging.
- leveldb, Imdb: database io for your program. > Need to know this for preparing your own data.
- protobuf: an efficient and flexible way to define data structure. > Need to know this for defining new layers.

Caffe Tutorial

- [Nets, Layers, and Blobs](#): the anatomy of a Caffe model.
- [Forward / Backward](#): the essential computations of layered compositional models.
- [Loss](#): the task to be learned is defined by the loss.
- [Solver](#): the solver coordinates model optimization.
- [Interfaces](#): command line, Python, and MATLAB Caffe.
- [Data](#): how to caffeinate data for model input.

<http://caffe.berkeleyvision.org/tutorial/>

Caffe

- Blob: Storage and Communication of Data
 - Data blobs are $N \times C \times H \times W$
- Net: Contains all the layers in the networks
 - Performs forward/backward pass through the entire network
- Solver: Used to set training/testing parameters
 - Number of iterations, back propagation method, etc..

Training: Step 1

- Create a lenet_train.prototxt
- Data Layers
- Operational Layers
- Loss Layers

Network Definition(train.prototxt)

```
name: "LeNet"
layer {
    name: "mnist"
    type: "Data"
    top: "data"
    top: "label"
    include {
        phase: TRAIN
    }
    transform_param {
        scale: 0.00390625
    }
    data_param {
        source: "examples/mnist/mnist_train_lmdb"
        batch_size: 64
        backend: LMDB
    }
}
```

```
layer {
    name: "conv1"
    type: "Convolution"
    bottom: "data"
    top: "conv1"
    param {
        lr_mult: 1
    }
    param {
        lr_mult: 2
    }
    convolution_param {
        num_output: 20
        kernel_size: 5
        stride: 1
        weight_filler {
            type: "xavier"
        }
        bias_filler {
            type: "constant"
        }
    }
}
```

Network Definition(train.prototxt)

```
layer {
    name: "pool2"
    type: "Pooling"
    bottom: "conv2"
    top: "pool2"
    pooling_param {
        pool: MAX
        kernel_size: 2
        stride: 2
    }
}
```

```
layer {
    name: "relu1"
    type: "ReLU"
    bottom: "ip1"
    top: "ip1"
}
```

```
layer {
    name: "ip1"
    type: "InnerProduct"
    bottom: "pool2"
    top: "ip1"
    param {
        lr_mult: 1
    }
    param {
        lr_mult: 2
    }
    inner_product_param {
        num_output: 500
        weight_filler {
            type: "xavier"
        }
        bias_filler {
            type: "constant"
        }
    }
}
```

Network Definition(train.prototxt)

```
layer {
    name: "loss"
    type: "SoftmaxWithLoss"
    bottom: "ip2"
    bottom: "label"
    top: "loss"
}
```

Training: Step 2

- Create a lenet_solver.prototxt

```
train_net: "lenet_train.prototxt"
base_lr: 0.01
momentum: 0.9
weight_decay: 0.0005
max_iter: 10000
snapshot_prefix: "lenet_snapshot"
# ... and some other options ...
```

Solver(solver.prototxt)

```
# The [train]/test net protocol buffer definition
net: "examples/mnist/lenet_train_test.prototxt"
# test_iter specifies how many forward passes the test should carry out.
# In the case of MNIST, we have test batch size 100 and 100 test iterations,
# covering the full 10,000 testing images.
test_iter: 100
# Carry out testing every 500 [training] iterations.
test_interval: 500
# The base learning rate, momentum and the weight decay of the network.
base_lr: 0.01
momentum: 0.9
weight_decay: 0.0005
# The learning rate policy
lr_policy: "step"
gamma: 0.1
stepsize: 3000
# Display every 100 iterations
display: 100
# The maximum number of iterations
max_iter: 10000
# snapshot intermediate results
snapshot: 5000
snapshot_prefix: "examples/mnist/lenet"
# solver mode: CPU or GPU
solver_mode: GPU
```

Training: Step 2

- Some details on SGD parameters

$$V_{t+1} = \mu V_t - \alpha (\nabla L(W_t) + \lambda W_t)$$
$$W_{t+1} = W_t + V_{t+1}$$

Diagram illustrating the SGD update steps:

- Momentum: Points to the term μV_t .
- LR: Points to the term $\alpha (\nabla L(W_t) + \lambda W_t)$.
- Decay: Points to the term λW_t .

Training: Step 3

- # train LeNet
 - `caffe train -solver examples/mnist/lenet_solver.prototxt`
- # train on GPU 2
 - `caffe train -solver examples/mnist/lenet_solver.prototxt -gpu 2`
- # resume training from the half-way point snapshot
 - `caffe train -solver examples/mnist/lenet_solver.prototxt -snapshot examples/mnist/lenet_iter_5000.solverstate`

Network Definition(test.prototxt)

Previously

```
name: "LeNet"
layer {
  name: "mnist"
  type: "Data"
  top: "data"
  top: "label"
  include {
    phase: TRAIN
  }
  transform_param {
    scale: 0.00390625
  }
  data_param {
    source: "examples/mnist/mnist_train_lmdb"
    batch_size: 64
    backend: LMDB
  }
}
```

```
layer {
  name: "mnist"
  type: "Data"
  top: "data"
  top: "label"
  include {
    phase: TEST
  }
  transform_param {
    scale: 0.00390625
  }
  data_param {
    source: "examples/mnist/mnist_test_lmdb"
    batch_size: 100
    backend: LMDB
  }
}
```

Network Definition(test.prototxt)

Previously

```
layer {
  name: "loss"
  type: "SoftmaxWithLoss"
  bottom: "ip2"
  bottom: "label"
  top: "loss"
}
```

```
layer {
  name: "accuracy"
  type: "Accuracy"
  bottom: "ip2"
  bottom: "label"
  top: "accuracy"
  include {
    phase: TEST
  }
}
```

PyCaffe (Training in Python)

- Add caffe python directory to path and import caffe

```
caffe_root = '../' # this file should be run from {caffe_root}/examples (otherwise change this line)

import sys
sys.path.insert(0, caffe_root + 'python')
import caffe
```

Layers

```
from caffe import layers as L, params as P

def lenet(lmdb, batch_size):
    # our version of LeNet: a series of linear and simple nonlinear transformations
    n = caffe.NetSpec()

    n.data, n.label = L.Data(batch_size=batch_size, backend=P.Data.LMDB, source=lmdb,
                           transform_param=dict(scale=1./255), ntop=2)

    n.conv1 = L.Convolution(n.data, kernel_size=5, num_output=20, weight_filler=dict(type='xavier'))
    n.pool1 = L.Pooling(n.conv1, kernel_size=2, stride=2, pool=P.Pooling.MAX)
    n.conv2 = L.Convolution(n.pool1, kernel_size=5, num_output=50, weight_filler=dict(type='xavier'))
    n.pool2 = L.Pooling(n.conv2, kernel_size=2, stride=2, pool=P.Pooling.MAX)
    n.fcl = L.InnerProduct(n.pool2, num_output=500, weight_filler=dict(type='xavier'))
    n.relu1 = L.ReLU(n.fcl, in_place=True)
    n.score = L.InnerProduct(n.relu1, num_output=10, weight_filler=dict(type='xavier'))
    n.loss = L.SoftmaxWithLoss(n.score, n.label)

    return n.to_proto()
```

Define solver and train network

```
caffe.set_device(0)
caffe.set_mode_gpu()

### load the solver and create train and test nets
solver = None # ignore this workaround for lmdb data (can't instantiate two solvers on the same data)
solver = caffe.SGDSolver('mnist/lenet_auto_solver.prototxt')
```

```
solver.net.forward() # train net
```

Access Net data

```
# we use a little trick to tile the first eight images
imshow(solver.net.blobs['data'].data[:8, 0].transpose(1, 0, 2).reshape(28, 8*28), cmap='gray'); ax
is('off')
print 'train labels:', solver.net.blobs['label'].data[:8]

train labels: [ 5.  0.  4.  1.  9.  2.  1.  3.]
```



PyCaffe (Testing in Python)

```
#load the model
net = caffe.Net('models/bvlc_reference_caffenet/train.prototxt',
                 'models/bvlc_reference_caffenet/train_iter30000.caffemodel',
                 caffe.TEST)

# load input and configure preprocessing
transformer = caffe.io.Transformer({'data': net.blobs['data'].data.shape})
transformer.set_mean('data', np.load('ilsvrc_2012_mean.npy').mean(1).mean(1))
transformer.set_transpose('data', (2,0,1))
transformer.set_channel_swap('data', (2,1,0))
transformer.set_raw_scale('data', 255.0)

#note we can change the batch size on-the-fly
#since we classify only one image, we change batch size from 10 to 1
net.blobs['data'].reshape(1,3,227,227)

#load the image in the data layer
im = caffe.io.load_image('examples/images/cat.jpg')
net.blobs['data'].data[...] = transformer.preprocess('data', im)

#compute
out = net.forward()

# other possibility : out = net.forward_all(data=np.asarray([transformer.preprocess('data', im)]))

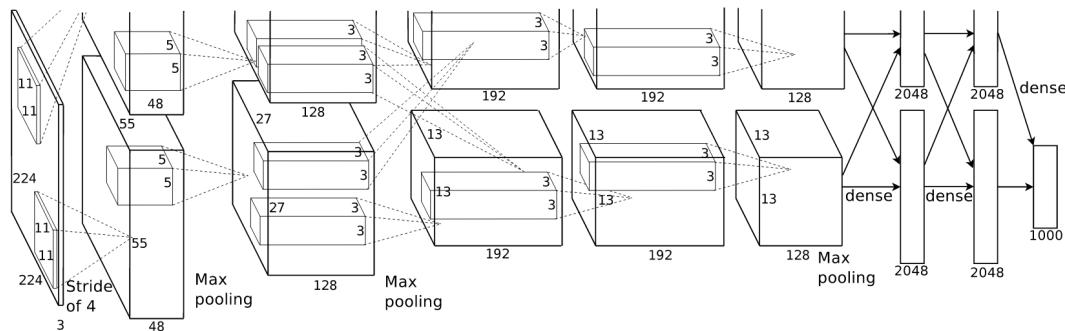
#predicted predicted class
print out['prob'].argmax()
```

Open Model Collection

- The Caffe Model Zoo
- open collection of deep models to share innovation
 - VGG ILSVRC14
 - Network-in-Network
 - MIT Places scene recognition model in the zoo
 - Help reproduce research
 - Bundled tools for loading and publishing models
- Share Your Models!

Reference Models

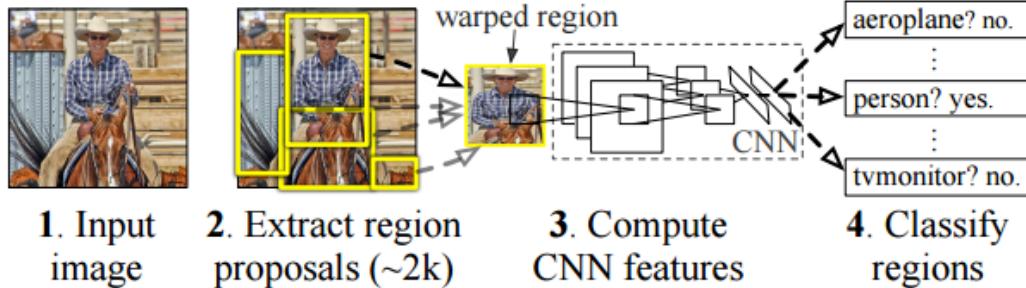
Alexnet: Imagenet Classification



Caffe offers the

- Model definitions
- Optimization settings
- Pre-trained weights so you can start right away.

R-CNN: *Regions with CNN features*



Why Fine-tuning?

- **Fine tuning** is a process to take a network model that has already been trained for a given task, and retrain it to make it perform a second similar task.
- Why?
 - More robust optimization
 - good initialization helps
 - Needs less data
 - Faster learning

Fine-tuning Tricks

- Learn the last layer first
 - Caffe layers have local learning rates: blobs_lr
 - Freeze all but the last layer for fast optimization and avoiding early divergence.
 - Stop if good enough, or keep fine-tuning
- Reduce the learning rate
 - Drop the solver learning rate by 10x, 100x –
 - Preserve the initialization from pre-training and avoid thrashing

Training tips

- Before running final/long training
 - Make sure you can overfit on a small training set
 - Make sure your loss decreases over first several iterations
 - Otherwise adjust parameter until it does, especially learning rate
- Separate train/val/test data



Any Questions?

